

# Linver Manual

July 4, 2016

Linver is a reference Matlab implementation of a verification framework for Bayesian inference algorithms. It is based on a linear regression problem for which analytical or semi-analytical solutions are known. As Markov chain Monte Carlo (MCMC) algorithms are a popular means of performing Bayesian inference, it provides a rigorous means of verifying the output chains of such algorithms are distributed correctly via an implementation of a hypothesis test for equal distributions based on the energy distance statistic (ADD CITATION). While the main goal of the code is as a reference for those interested in implementing the framework in other verification software, it is also useable as-is as a basic verification tool.

# Chapter 1

## Parameters of the Linear Equation

The verification framework is based on Bayesian inference for the linear equation defined by

$$y = G\beta + \varepsilon(\lambda, \phi).$$

The variables making up this equation are:

$y$  : A vector of  $N$  observations.

$G$  : The  $N \times N_\beta$  design matrix.

$\beta$  : The  $N_\beta \times 1$  column vector  $[\beta_1 \ \cdots \ \beta_{N_\beta}]^T$  of regression parameters.

$\varepsilon$  : The  $N \times 1$  column vector of observation noise.

$\lambda$  : The scale parameter for the noise.

$\phi$  : The correlation parameter for the noise covariance.

In this implementation of the framework, the first column of the design matrix  $G$  is all one's, so the first regression parameter  $\beta_1$  is a bias term. The remaining entries of  $G$  (if any) are drawn from a standard normal distribution (i.e., mean 0 and variance 1). This choice of  $G$  is not crucial to the framework as-implemented and can be substituted.

The observation noise is zero-mean and normally distributed so that  $\varepsilon \sim N(0, C)$  with the  $N \times N$  covariance matrix  $C(\lambda, \phi)$  depending on the parameters  $\lambda, \phi$ . The  $\lambda, \phi$  dependence is expressed by

$$C(\lambda, \phi) = \frac{1}{\lambda} R(\phi),$$

where  $R(\phi)$  is the correlation function which depends on  $\phi$ . The framework includes three possible choices

of correlation:

$$\begin{aligned}
\text{No correlation: } R(\phi) &= \begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix} \\
\text{Equicorrelation: } R(\phi) &= \begin{bmatrix} 1 & \phi & \phi & \cdots & \phi & \phi \\ \phi & 1 & \phi & \cdots & \phi & \phi \\ \phi & \phi & 1 & \cdots & \phi & \phi \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ \phi & \phi & \phi & \cdots & 1 & \phi \\ \phi & \phi & \phi & \cdots & \phi & 1 \end{bmatrix} \\
\text{AR(1) correlation: } R(\phi) &= \begin{bmatrix} 1 & \phi & \phi^2 & \cdots & \phi^{N-2} & \phi^{N-1} \\ \phi & 1 & \phi & \cdots & \phi^{N-3} & \phi^{N-2} \\ \phi^2 & \phi & 1 & \cdots & \phi^{N-4} & \phi^{N-3} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ \phi^{N-2} & \phi^{N-3} & \phi^{N-4} & \cdots & 1 & \phi \\ \phi^{N-1} & \phi^{N-2} & \phi^{N-3} & \cdots & \phi & 1 \end{bmatrix}
\end{aligned}$$

For each of the correlation cases, the correlation parameter  $\phi$  is restricted to a different domain:

$$\begin{aligned}
\text{No correlation: } \phi &\in \emptyset \\
\text{Equicorrelation: } \phi &\in [0, 1) \\
\text{AR(1) correlation: } \phi &\in (-1, 1)
\end{aligned}$$

Note that although we have defined the domain for  $\phi$  in the no correlation case to be the empty set since the corresponding correlation matrix in this case has no  $\phi$  dependence. The reason for using these choices of covariance functions is that these choices each have known analytical forms for the inverse and determinant. This facilitates computation of the exact solution to the Bayesian inverse problems described below, as the inverse and determinant of the covariance matrices appear in the analytical and semi-analytical solutions for the parameter distributions.

## Chapter 2

# Bayesian Inference

The purpose of the framework is to provide a means of verifying Bayesian inference algorithms using the equations and parameters in Chapter 1. The framework has three different cases of unknown parameters depending on which of the parameters  $\beta$ ,  $\lambda$ , and  $\phi$  are treated as unknown. The cases used in the framework are summarized in the following table.

Case	Known	Unknown	Calibration Parameters
1	$\lambda, \phi$	$\beta$	$\hat{\beta}$
2	$\phi$	$\beta, \lambda$	$\hat{\beta}, \hat{\lambda}$
3	None	$\beta, \lambda, \phi$	$\hat{\beta}, \hat{\lambda}, \hat{\phi}$

Bayesian inference relies on Bayes rule

$$p(\hat{\theta}|y) = \frac{p(y|\hat{\theta})p_0(\hat{\theta})}{p(y)},$$

where  $\hat{\theta}$  is the inferred parameter,  $y$  is the data, and the  $p(\cdot)$  functions refer to the following likelihoods:

$p(\hat{\theta}|y)$  : Posterior of parameter value  $\hat{\theta}$  given data  $y$

$p(y|\hat{\theta})$  : Likelihood of data  $y$  given parameter value  $\hat{\theta}$

$p_0(\hat{\theta})$  : Prior for inferred parameter value  $\hat{\theta}$

$p(y)$  : Marginal likelihood of the data  $y$  given by  $\int_{D(\hat{\theta})} p(y|\hat{\theta})p_0(\hat{\theta})d\hat{\theta}$

As can be determined from the above table and Bayes rule above, the verification framework computes the following marginal posteriors according to which set of parameters are chosen to be calibrated.

**Case 1** :  $p(\hat{\beta}|y)$

**Case 2** :  $p(\hat{\beta}|y), p(\hat{\lambda}|y)$

**Case 3** :  $p(\hat{\beta}|y), p(\hat{\lambda}|y), p(\hat{\phi}|y)$

Case 1 and Case 2 are computed exactly, in the sense that an exact parametric distribution for each of the marginal posteriors can be derived which can be computed using standard routines for calculating the

associated special functions. The posterior for Case 3 is an integral expression involving special functions and is computed using numerical integration. Note that Case 3 requires the use of the equicorrelation or AR(1) correlation function, since there is no  $\phi$  dependence for uncorrelated observation error.

For each of the three cases of unknown parameters, there are two choices of prior. These are labeled the “Non-informative” and “Gaussian” priors, according to whether the marginal prior for  $\hat{\beta}$  is uniform or Gaussian. The following describes the priors for each of the calibrated variables.

### Non-informative

$$\text{Case 1 : } p_0(\hat{\beta}) \propto 1$$

$$\text{Case 2 : } p_0(\hat{\beta} | \hat{\lambda}) \propto 1$$

$$p_0(\hat{\lambda}) \propto \frac{1}{\hat{\lambda}}$$

$$\text{Case 3 : } p_0(\hat{\beta} | \hat{\lambda}, \hat{\phi}) \propto 1$$

$$p_0(\hat{\lambda} | \hat{\phi}) \propto \frac{1}{\hat{\lambda}}$$

$$p_0(\hat{\phi}) = \frac{1}{\phi_H - \phi_L}$$

The priors which are described as a proportional expression are improper priors. The parameters  $\phi_H$  and  $\phi_L$  are the high and low values of the domain of  $\phi$ , so that  $\phi \in [\phi_L, \phi_H]$ . These are configurable parameters in the framework and will depend on the choice of correlation function. These are the only parameters needed for the non-informative prior, and only when  $\phi$  is an unknown parameter (i.e., only in Case 3).

### Gaussian

$$\text{Case 1 : } p_0(\hat{\beta}) \text{ is } N\left(\mu_0, \frac{1}{\lambda}\Sigma_0\right)$$

$$\text{Case 2 : } p_0(\hat{\beta} | \hat{\lambda}) \text{ is } N\left(\mu_0, \frac{1}{\hat{\lambda}}\Sigma_0\right)$$

$$p_0(\hat{\lambda}) \propto \frac{1}{\hat{\lambda}}$$

$$\text{Case 3 : } p_0(\hat{\beta} | \hat{\lambda}, \hat{\phi}) \text{ is } N\left(\mu_0, \frac{1}{\hat{\lambda}}\Sigma_0\right)$$

$$p_0(\hat{\lambda} | \hat{\phi}) \propto \frac{1}{\hat{\lambda}}$$

$$p_0(\hat{\phi}) = \frac{1}{\phi_H - \phi_L}$$

The only difference between the non-informative and Gaussian cases is the prior on  $\hat{\beta}$  is Gaussian with a prescribed  $N_\beta \times 1$  mean vector  $\mu_0$  and unscaled  $N_\beta \times N_\beta$  covariance  $\Sigma_0$ . The framework uses diagonal  $\Sigma_0$ , although this is not strictly necessary. Make note that the prior distribution in Case 1 has a covariance matrix scaled by  $1/\lambda$ , the known value of  $\lambda$ . In the other cases  $\lambda$  is unknown so this scaling is replaced by  $1/\hat{\lambda}$ , the calibrated variable.

## Chapter 3

# Using the Verification Framework

The following is an outline of the steps involved in using the verification framework.

1. Choose the size of the problem:  $N$ , the number of observations and  $N_\beta$ , the number of regression parameters.
2. Choose the values of the true parameters  $\beta, \lambda, \phi$ . These are the values that the Bayesian inference algorithm will attempt to learn from the data.
3. Set the design matrix  $G$ . The default recommendation (and what is used in examples) is to set all entries in the first column to ones and fill the remaining entries (if any) with data drawn from a standard normal distribution.
4. Use the true parameter  $\beta$  to determine the error-free model output  $y_0$  according to  $y_0 = G\beta$ .
5. Choose a correlation function  $R(\phi)$ : no correlation, equicorrelation, or AR(1) correlation.
6. Use the chosen correlation function and the true parameters  $\lambda, \phi$  to generate the observation error  $\varepsilon$  according to the distribution  $N(0, C)$  where  $C = R(\phi)/\lambda$ .
7. Add the observation error  $\varepsilon$  to the error-free model output  $y_0$  to obtain the calibration data  $y = y_0 + \varepsilon$ .
8. Choose which parameters to calibrate (i.e., which are treated as unknowns): Case 1:  $\hat{\beta}$ , Case 2:  $\hat{\beta}, \hat{\lambda}$ , or Case 3:  $\hat{\beta}, \hat{\lambda}, \hat{\phi}$ . If using Case 3, choose the parameters  $\phi_L, \phi_H$  defining restricting the domain of  $\phi$ . These need to be chosen in agreement with the domain associated with the correlation function.
9. Choose whether to use the non-informative prior or the Gaussian prior for the regression parameters  $\hat{\beta}$ . If using the Gaussian prior, choose the prior mean  $\mu_0$  and prior un-scaled covariance  $\Sigma_0$ .
10. Using the calibration routine to be verified, calculate the posterior for the unknown parameters. It is important that this be done using the calibration data  $y$  generated above, using the same prior.
11. Use the verification framework to calculate the true posterior using all the chosen parameters as input.
12. Compare the true posterior computed by the verification framework with that computed by the routine being verified.

This outline leaves open the specific ways the comparison between the true posterior and the output of the calibration routine are to be done. There are multiple ways to do this, some being more appropriate for different types of calibration routines. We specifically mention the case of verifying Markov chain Monte Carlo (MCMC) output, as we have included a method aimed at this common case. Here we recommend a statistical hypothesis test based on the energy statistic, and the Matlab implementation includes an implementation of this.





## Chapter 4

# Using the Matlab Implementation

The Matlab scripts `demo_case1.m`, `demo_case2.m`, `demo_case3.m` demonstrate the use of the Matlab implementation for Case 1 ( $\beta$  unknown), Case 2 ( $\beta, \lambda$  unknown), and Case 3 ( $\beta, \lambda, \phi$  unknown) respectively.

The main function used in computing the posterior is `eval_posterior`. This function takes a structure `param` as input which defines all the parameters of the problem. The fields of this structure are as follows.

<code>param.N:</code>	Number of observations.
<code>param.Nbeta:</code>	Number of regression parameters.
<code>param.G:</code>	N by Nbeta design matrix.
<code>param.prior:</code>	Structure defining the prior. See below for more information.
<code>param.beta:</code>	Nbeta by 1 vector of the true regression parameters $\beta$
<code>param.lambda:</code>	Scalar value of the true scale parameter $\lambda$ .
<code>param.phi:</code>	Scalar value of the true correlation parameter $\phi$ .
<code>param.corrfunc:</code>	String defining the correlation function type: "none", "equal", "ar".
<code>param.y:</code>	N by 1 vector of observations.
<code>param.unknowns:</code>	String defining which parameters are unknown. See below.
<code>param.betarange:</code>	Nbeta by 2 matrix. Left column is lower limit of $\beta$ , right column is upper limit.
<code>param.lambdarange:</code>	1 by 2 matrix. Left value is lower limit of $\lambda$ , right column is upper limit.
<code>param.phirange:</code>	1 by 2 matrix. Left value is lower limit of $\phi$ , right column is upper limit.

**Prior structure.** Let the prior structure be named `prior`. It always has the field `prior.type` which is a string variable with two possible values: "noninformative" or `gaussian`. If the `prior.type` field is set to "noninformative", no other fields need to be defined. If it is set to `gaussian`, the field `prior.sigma0` must be the Nbeta by Nbeta unscaled covariance and `prior.mu0` is the Nbeta by 1 mean vector of the prior.

**Unknowns strings.** The `prior.unknowns` field can take one of the string values: "beta", "beta\_lambda", "beta\_lambda\_phi". These correspond to whether the unknowns are Case 1:  $\beta$ , Case 2:  $\beta, \lambda$ , or Case 3:  $\beta, \lambda, \phi$ .

**Calibration data.** The function `eval_noise` takes the same `param` structure above as input, except `param.y` need not be defined. It returns a N by 1 vector of unscaled observation error. This vector can be divided by  $\sqrt{\lambda}$  and added to the error-free model output to provide calibration data.

**Posterior structure.** The function `eval_posterior` returns a structure which we call `post`. The fields of this structure depend on which parameters are unknown.

NEXT: Describe what the posterior structure has in each of the unknown cases.

NEXT: Describe the use of the energy distance test functions.